

```
In [2]: import pandas as pd
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.datasets import load_iris
```

```
In [3]: iris=load_iris()
X=iris.data
y=iris.target
X_train,X_test,y_train,y_test=train_test_split(X,y,random_state=1)
```

```
In [5]: X.shape
print(iris.DESCR)
```

```
.. _iris_dataset:
```

```
Iris plants dataset
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 150 (50 in each of three classes)
:Number of Attributes: 4 numeric, predictive attributes and the class
:Attribute Information:
  - sepal length in cm
  - sepal width in cm
  - petal length in cm
  - petal width in cm
  - class:
    - Iris-Setosa
    - Iris-Versicolour
    - Iris-Virginica
```

```
:Summary Statistics:
```

	Min	Max	Mean	SD	Class Correlation
sepal length:	4.3	7.9	5.84	0.83	0.7826
sepal width:	2.0	4.4	3.05	0.43	-0.4194
petal length:	1.0	6.9	3.76	1.76	0.9490 (high!)
petal width:	0.1	2.5	1.20	0.76	0.9565 (high!)

```
:Missing Attribute Values: None
:Class Distribution: 33.3% for each of 3 classes.
:Creator: R.A. Fisher
:Donor: Michael Marshall (MARSHALL%PLU@io.arc.nasa.gov)
:Date: July, 1988
```

The famous Iris database, first used by Sir R.A. Fisher. The dataset is taken from Fisher's paper. Note that it's the same as in R, but not as in the UCI Machine Learning Repository, which has two wrong data points.

This is perhaps the best known database to be found in the pattern recognition literature. Fisher's paper is a classic in the field and is referenced frequently to this day. (See Duda & Hart, for example.) The data set contains 3 classes of 50 instances each, where each class refers to a type of iris plant. One class is linearly separable from the other 2; the latter are NOT linearly separable from each other.

```
.. topic:: References
```



<b>5073</b>	169142	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850	0.697211	-2.064945	...
<b>5074</b>	169347	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170	0.248525	-1.127396	...
<b>5075</b>	169351	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739	1.210158	-0.652250	...
<b>5076</b>	169966	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002	1.058733	-1.632333	...
<b>5077</b>	170348	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050	-0.068384	0.577829	...

5078 rows × 31 columns

In [29]: `df.Class.value_counts()`

Out[29]:

```

0      4586
1       492
Name: Class, dtype: int64

```

In [30]: `X=df.iloc[:,1:-2].values`  
`y=df.iloc[:, -1].values`

In [31]: `df.describe()`

Out[31]:

	Time	V1	V2	V3	V4	V5	V6	V7
<b>count</b>	5078.000000	5078.000000	5078.000000	5078.000000	5078.000000	5078.000000	5078.000000	5078.000000
<b>mean</b>	26382.185112	-0.705344	0.584418	-0.103781	0.425607	-0.309467	-0.108354	-0.414131
<b>std</b>	52274.968414	2.863990	2.026708	3.356264	2.102100	2.234907	1.412898	2.988556
<b>min</b>	0.000000	-30.552380	-15.732974	-31.103685	-4.657545	-32.092129	-7.465603	-43.557242
<b>25%</b>	978.000000	-1.179498	-0.167130	-0.265157	-0.800510	-0.678437	-0.838720	-0.490238
<b>50%</b>	2098.500000	-0.462149	0.447474	0.628581	0.224502	-0.106695	-0.292258	0.087072
<b>75%</b>	3321.750000	1.082291	1.037522	1.291780	1.185083	0.433680	0.374438	0.603359
<b>max</b>	170348.000000	2.355634	22.057729	4.017561	12.114672	11.095089	21.393069	34.303177

8 rows × 31 columns

In [32]: `X_train,X_test,y_train,y_test=train_test_split(X,y,random_state=1)`  
`model=KNeighborsClassifier()`  
`model.fit(X_train,y_train)`  
`print(model.score(X_train,y_train))`  
`print(model.score(X_test,y_test))`

```

0.9855567226890757
0.9921259842519685

```

In [33]: `from sklearn.model_selection import cross_val_score`

In [34]: `model=KNeighborsClassifier()`  
`cross_val_score(model,X,y,cv=5).mean()`

Out[34]: 0.9836552112020479

**accuracy score---->overall performance of model  
by considering all classes**

# f1 score----->performance of model with class 1

```
In [35]: from sklearn.metrics import f1_score
```

```
In [37]: model=KNeighborsClassifier()
model.fit(X_train,y_train)
pred_train=model.predict(X_train)
pred_test=model.predict(X_test)
print("Train F1:",f1_score(y_train,pred_train))
print("Test F1:",f1_score(y_test,pred_test))
```

Train F1: 0.9237170596393898  
Test F1: 0.9523809523809523

```
In [39]: X_train,X_test,y_train,y_test=train_test_split(X,y,random_state=1)
sc=StandardScaler()
X_train_new=sc.fit_transform(X_train)
X_test_new=sc.transform(X_test)
model=KNeighborsClassifier()
model.fit(X_train_new,y_train)
pred_train=model.predict(X_train_new)
pred_test=model.predict(X_test_new)
print("Train F1:",f1_score(y_train,pred_train))
print("Test F1:",f1_score(y_test,pred_test))
```

Train F1: 0.9103641456582633  
Test F1: 0.9326923076923077

```
In [40]: print("Train F1:",f1_score(y_train,pred_train,average=None))
print("Test F1:",f1_score(y_test,pred_test,average=None))
```

Train F1: [0.99072733 0.91036415]  
Test F1: [0.99399657 0.93269231]

```
In [41]: #SMOTE
```

```
Cell In[41], line 2
    pip install imblearn
      ^
```

**SyntaxError:** invalid syntax

```
In [42]: pip install imblearn
```

Collecting imblearn  
 Downloading imblearn-0.0-py2.py3-none-any.whl (1.9 kB)  
Requirement already satisfied: imbalanced-learn in c:\users\ducat\anaconda3\lib\site-packages (from imblearn) (0.10.1)  
Requirement already satisfied: numpy>=1.17.3 in c:\users\ducat\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.24.3)  
Requirement already satisfied: scipy>=1.3.2 in c:\users\ducat\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.10.1)  
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\ducat\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.2.2)  
Requirement already satisfied: joblib>=1.1.1 in c:\users\ducat\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.2.0)  
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\ducat\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (2.2.0)  
Installing collected packages: imblearn  
Successfully installed imblearn-0.0  
Note: you may need to restart the kernel to use updated packages.

```
In [43]: from imblearn.over_sampling import SMOTE
```

```
In [47]: sm=SMOTE()  
X1,y1=sm.fit_resample(X,y)
```

```
In [48]: X.shape
```

```
Out[48]: (5078, 28)
```

```
In [49]: X1.shape
```

```
Out[49]: (9172, 28)
```

```
In [50]: X_train,X_test,y_train,y_test=train_test_split(X1,y1,random_state=1)  
sc=StandardScaler()  
X_train_new=sc.fit_transform(X_train)  
X_test_new=sc.transform(X_test)  
model=KNeighborsClassifier()  
model.fit(X_train_new,y_train)  
pred_train=model.predict(X_train_new)  
pred_test=model.predict(X_test_new)  
print("Train F1:",f1_score(y_train,pred_train,average=None))  
print("Test F1:",f1_score(y_test,pred_test,average=None))
```

```
Train F1: [0.98699021 0.98713315]
```

```
Test F1: [0.9845338 0.98493328]
```

```
In [ ]:
```